

18IT140_PRACTICAL_12_House_Price_Prediction_with_Python

October 29, 2021

1 Machine Learning Project on House Price Prediction with Python

```
[1]: import pandas as pd
housing = pd.read_csv("/content/sample_data/housing.csv")
housing.head()
```

```
[1]:  longitude  latitude  ...  median_house_value  ocean_proximity
0    -122.23    37.88  ...           452600.0          NEAR BAY
1    -122.22    37.86  ...           358500.0          NEAR BAY
2    -122.24    37.85  ...           352100.0          NEAR BAY
3    -122.25    37.85  ...           341300.0          NEAR BAY
4    -122.25    37.85  ...           342200.0          NEAR BAY
```

[5 rows x 10 columns]

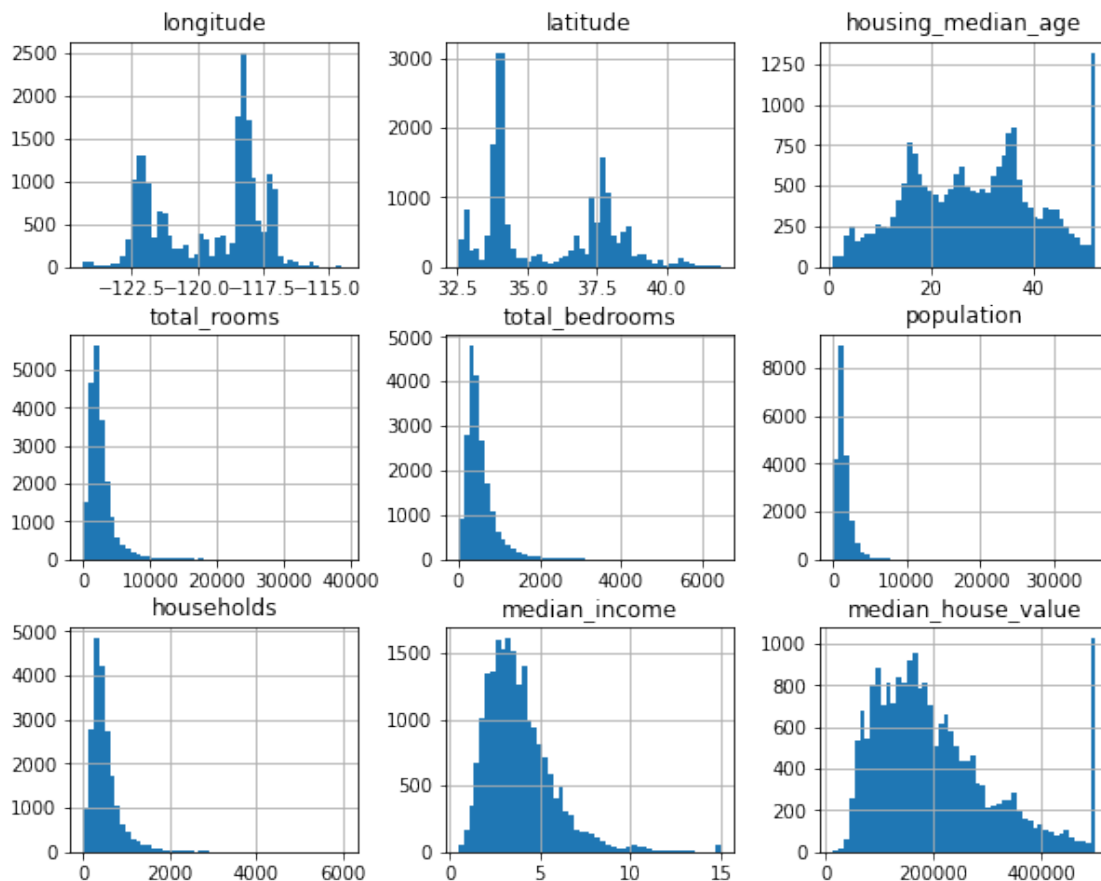
```
[2]: housing.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude             20640 non-null  float64
1   latitude              20640 non-null  float64
2   housing_median_age    20640 non-null  float64
3   total_rooms           20640 non-null  float64
4   total_bedrooms        20433 non-null  float64
5   population            20640 non-null  float64
6   households            20640 non-null  float64
7   median_income         20640 non-null  float64
8   median_house_value    20640 non-null  float64
9   ocean_proximity       20640 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

```
[3]: housing.ocean_proximity.value_counts()
```


```
[3]: <1H OCEAN      9136  
     INLAND       6551  
     NEAR OCEAN   2658  
     NEAR BAY     2290  
     ISLAND        5  
     Name: ocean_proximity, dtype: int64
```

```
[4]: import matplotlib.pyplot as plt  
     housing.hist(bins=50, figsize=(10, 8))  
     plt.show()
```



```
[5]: from sklearn.model_selection import train_test_split  
     train_set, test_set = train_test_split(housing, test_size=0.2, random_state=42)
```

```
[6]: import numpy as np  
     housing['income_cat'] = pd.cut(housing['median_income'], bins=[0., 1.5, 3.0, 4.  
     →5, 6., np.inf], labels=[1, 2, 3, 4, 5])  
     housing['income_cat'].hist()  
     plt.show()
```




```
[7]: from sklearn.model_selection import StratifiedShuffleSplit
split = StratifiedShuffleSplit(n_splits=1, test_size=0.2, random_state=42)
for train_index, test_index in split.split(housing, housing["income_cat"]):
    strat_train_set = housing.loc[train_index]
    strat_test_set = housing.loc[test_index]
print(strat_test_set['income_cat'].value_counts() / len(strat_test_set))
```

```
3    0.350533
2    0.318798
4    0.176357
5    0.114583
1    0.039729
```

Name: income_cat, dtype: float64

```
[8]: for set_ in (strat_train_set, strat_test_set):
    set_.drop('income_cat', axis=1, inplace=True)
housing = strat_train_set.copy()
```

```
[9]: housing.plot(kind='scatter', x='longitude', y='latitude', alpha=0.4,
    →s=housing['population']/100, label='population',
figsize=(12, 8), c='median_house_value', cmap=plt.get_cmap('jet'),
    →colorbar=True)
plt.legend()
plt.show()
```



```
[10]: corr_matrix = housing.corr()
print(corr_matrix.median_house_value.sort_values(ascending=False))
```

```
median_house_value    1.000000
median_income         0.687160
total_rooms           0.135097
housing_median_age    0.114110
households            0.064506
total_bedrooms        0.047689
population            -0.026920
longitude             -0.047432
latitude              -0.142724
Name: median_house_value, dtype: float64
```

```
[11]: housing["rooms_per_household"] = housing["total_rooms"]/housing["households"]
housing["bedrooms_per_room"] = housing["total_bedrooms"]/housing["total_rooms"]
housing["population_per_household"] = housing["population"]/
    ↪ housing["households"]

corr_matrix = housing.corr()
print(corr_matrix["median_house_value"].sort_values(ascending=False))
```

```
median_house_value    1.000000
median_income         0.687160
```

```
rooms_per_household      0.146285
total_rooms               0.135097
housing_median_age       0.114110
households                0.064506
total_bedrooms           0.047689
population_per_household -0.021985
population                -0.026920
longitude                 -0.047432
latitude                 -0.142724
bedrooms_per_room        -0.259984
Name: median_house_value, dtype: float64
```

```
[12]: # Data Preparation
housing = strat_train_set.drop("median_house_value", axis=1)
housing_labels = strat_train_set["median_house_value"].copy()

median = housing["total_bedrooms"].median()
housing["total_bedrooms"].fillna(median, inplace=True)

housing_num = housing.drop("ocean_proximity", axis=1)

from sklearn.base import BaseEstimator, TransformerMixin

# column index
rooms_ix, bedrooms_ix, population_ix, households_ix = 3, 4, 5, 6

class CombinedAttributesAdder(BaseEstimator, TransformerMixin):
    def __init__(self, add_bedrooms_per_room=True): # no *args or **kwargs
        self.add_bedrooms_per_room = add_bedrooms_per_room
    def fit(self, X, y=None):
        return self # nothing else to do
    def transform(self, X):
        rooms_per_household = X[:, rooms_ix] / X[:, households_ix]
        population_per_household = X[:, population_ix] / X[:, households_ix]
        if self.add_bedrooms_per_room:
            bedrooms_per_room = X[:, bedrooms_ix] / X[:, rooms_ix]
            return np.c_[X, rooms_per_household, population_per_household,
                          bedrooms_per_room]
        else:
            return np.c_[X, rooms_per_household, population_per_household]
```

```
[13]: from sklearn.preprocessing import OneHotEncoder
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer
num_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy="median")),
    ('attribs_adder', CombinedAttributesAdder()),
```

```

        ('std_scaler', StandardScaler()),
    ])
    housing_num_tr = num_pipeline.fit_transform(housing_num)

    from sklearn.compose import ColumnTransformer
    num_attribs = list(housing_num)
    cat_attribs = ["ocean_proximity"]
    full_pipeline = ColumnTransformer([
        ("num", num_pipeline, num_attribs),
        ("cat", OneHotEncoder(), cat_attribs),
    ])
    housing_prepared = full_pipeline.fit_transform(housing)

```

2 Linear Regression for House Price Prediction with Python

```

[14]: from sklearn.linear_model import LinearRegression
lin_reg = LinearRegression()
lin_reg.fit(housing_prepared, housing_labels)

data = housing.iloc[:5]
labels = housing_labels.iloc[:5]
data_preparation = full_pipeline.transform(data)
print("Predictions: ", lin_reg.predict(data_preparation))

```

```

Predictions: [210644.60459286 317768.80697211 210956.43331178  59218.98886849
 189747.55849879]

```

I hope you liked this article on Machine Learning project on House Price Prediction with Python.

```

[ ]: !wget -nc https://raw.githubusercontent.com/brpy/colab-pdf/master/colab_pdf.py
from colab_pdf import colab_pdf
colab_pdf('18IT140_PRACTICAL_12_House_Price_Prediction_with_Python.ipynb')

```

```

--2021-10-29 15:10:39-- https://raw.githubusercontent.com/brpy/colab-
pdf/master/colab_pdf.py
Resolving raw.githubusercontent.com (raw.githubusercontent.com)...
185.199.108.133, 185.199.110.133, 185.199.111.133, ...
Connecting to raw.githubusercontent.com
(raw.githubusercontent.com)|185.199.108.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 1864 (1.8K) [text/plain]
Saving to: colab_pdf.py

```

```

colab_pdf.py          100%[=====>]    1.82K  --.-KB/s    in 0s

```

2021-10-29 15:10:39 (40.1 MB/s) - colab_pdf.py saved [1864/1864]

WARNING: apt does not have a stable CLI interface. Use with caution in scripts.

WARNING: apt does not have a stable CLI interface. Use with caution in scripts.

Extracting templates from packages: 100%